



High Performance Computing (HPC)

- Introduction to HPC and why it's useful
- How do you parallelize your code?
- The practice of running software on a cluster



Introduction to HPC and why it's useful

Gordon Moore

Born January 3, 1929 (age 84)
San Francisco, California, USA


Nationality American

Alma mater University of California, Berkeley;
California Institute of Technology

Occupation Chairman Emeritus, Intel
Corporation

Net worth ▲ \$4 billion USD (2011)^[1]

Jan 2015 update: \$6.7 billion



Introduction to HPC and why it's useful


- Embarrassingly parallel problems
 - Graphics
 - Simulations with multiple parameters
- Non embarrassingly parallel problems
 - Fluid dynamics
 - A lot of the tasks run by a single program




How do you parallelize your code? !

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
– Should be used in your code to give a simulation number

Before



Now




Using your PC

```
for ( i in 1 : 10 )
{
  do_simulation( i )
}
```

Shell script on
the cluster


Using HPC




How do you parallelize your code? !

- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
 - Should be used in your code to give a simulation number

Before




Now




How do you parallelize your code? !


- `as.numeric(Sys.getenv("PBS_ARRAY_INDEX"))`
 - Should be used in your code to give a simulation number
- Pseudo random numbers
 - Given a certain random number seed, you get the same sequence of random numbers every time
 - Each simulation should have a different seed



Handling memory




Handling memory




Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically



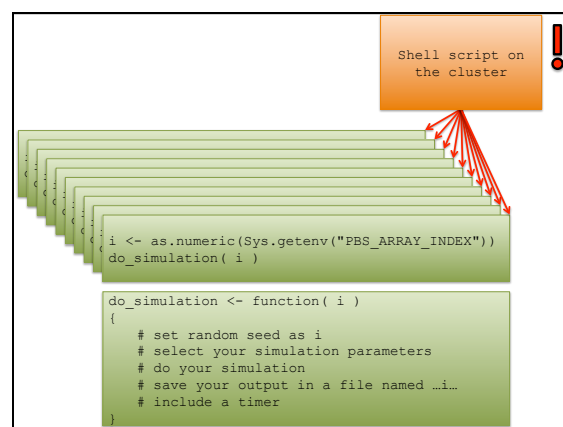
Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically

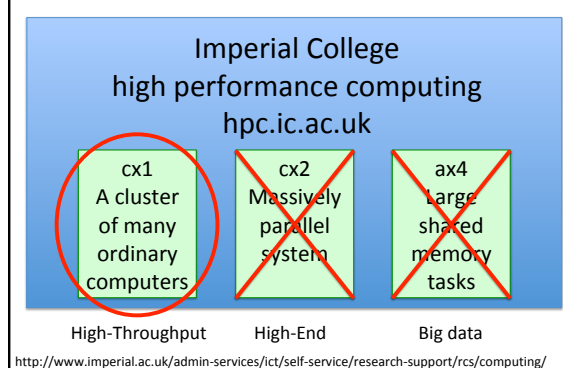


Handling memory

- Save your results in memory and then write to disk at the end
- Output your code to a series of files
- Write local code to read in your series of files automatically
- Build a timer into your code
- Test your code locally to know your memory and time requirements



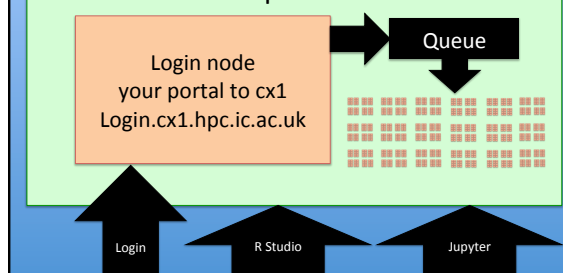
The practice of running code on a cluster



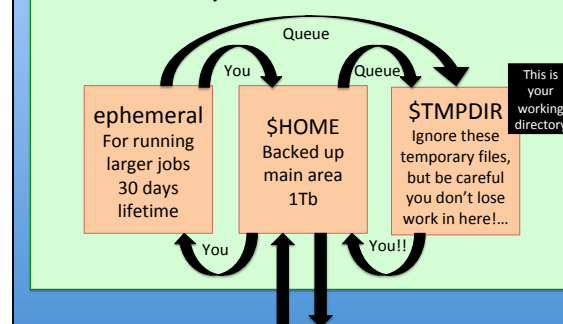
The Imperial College high performance computing PC cluster cx1.hpc.ic.ac.uk



The Imperial College high performance computing PC cluster cx1.hpc.ic.ac.uk



Where your data is stored...



Step 1: get your code onto the cluster

- Go to filezilla-project.org
- Download FileZilla Client and install it.
- Open FileZilla and put in the following settings
 - Host: sftp://login.cx1.hpc.ic.ac.uk
 - Username: your username for all Imperial services
 - Password: your usual Imperial password
 - Port: 22
- Press the Quickconnect button
- Copy your files and folders across onto HPC.

Step 1: get your code onto the cluster

Alternative method if you prefer command line...

- Use sftp: from the directory of your code in a shell window type
 - sftp `username@login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - put `filename.R`
 - exit
- Or use scp:
 - scp `path/to/file.txt`
`username@login.cx1.hpc.ic.ac.uk:/home/username/`
- Also see separate notes to be shared

Step 2: log into the cluster

- Use ssh: from a shell window type
 - Ssh -l `username` `login.cx1.hpc.ic.ac.uk`
 - You will be asked for your standard cluster password
 - Now it's as though you were sitting with a shell open at the login node.
 - ls (will list the files in \$HOME)
 - mkdir `foldername` (make a new folder)
 - mv `filename` `$HOME/foldername` (move)
 - cd `foldername` (change directory)
 - cat `filename` (see your file to check it's contents)
 - anaconda-setup (set up anaconda – one time only)
 - conda install `r` (install R – one time only)

Step 3: make a file for your shell script

- This is the list of instructions to be executed when you get to the front of the queue is written in shell script. It should be a .sh file
- **Do not** run code on the login node – always write a shell script and wait in the queue.
- If you type `cat > filename.sh`
- You will then get the chance to type text (pressing enter for new lines) and the cat command will make the file containing the text that you typed.
- When you are finished typing the contents of your new file press control and D to complete the process.
- Type `cat filename.sh` to check that your file is correct before submitting it to the queue.

Step 3 continued: your shell script file

```
#!/bin/bash
#PBS -l walltime=12:00:00
#PBS -l select=1:ncpus=1:mem=1gb
module load anaconda3/personal
echo "R is about to run"
R --vanilla < $HOME/Rtest/ForwardsNTC_V5.R
mv datafilename* $HOME
echo "R has finished running"
# this is a comment at the end of the file
```

You can run Python code too –
just use different commands here

Step 3 continued: job sizing

Job class	Number of nodes N	ncpus/node	Max mem/node	Max walltime/hr	Max number of running jobs per user
throughput	1	1-8	96GB	up to 72hr	unlimited for jobs <=24hr in length
general	1 - 16	32	62GB or 124GB	up to 72hr	unlimited for jobs <=24hr in length
singlenode	1	48	124GB	up to 24hr	10
multinode	3 - 16	12	46GB	up to 48hr	unlimited
debug	1	1-8	96GB	up to 30 mins	1
large memory	1	12 or 24	190 or 380GB	48 hr	unlimited
GPU	1-2	1-8	1 - 32GB	48hr 72hr (P1000 only)	8
long	1	1-8	96GB	72 - 1000 hr	1

Step 4: submitting your job to the cluster!

- You are now in the \$HOME directory with your code and shell script both written.
- To submit your job type

```
qsub -J 1-32 filename.sh
```

```
qstat
```

(S changes from Q to B when running)
- If you want to delete a job

```
qstat
```

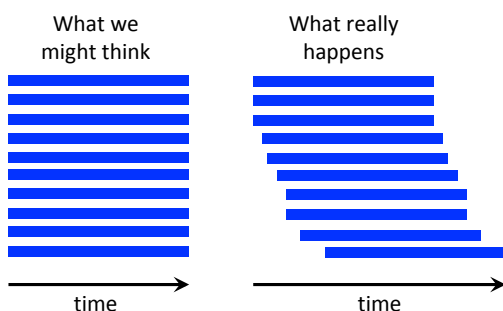
```
qdel job-id[]
```

(the [] is for array jobs only)
- qstat will give you a list of jobs and you would get the job-id from there.

Step 5: check that all is well

- Wait 5-10 minutes then check that nothing has gone wrong.
- qstat (is your job running still)
- ls (are output files as expected)
- cat filename.sh.ejob-id.index (are error files empty?)
- cat filename.sh.ojob-id.index (are standard output files as expected)
- qstat (is your job running still)
- exit (you're done for now come back later)

Step 5 continued: how jobs are run



Step 6: Getting your results back


- qstat (is your job running still)
- cd \$HOME
- ls (output files as expected?)
- cat output filename (contents as expected?)
- cat filename.sh.ejob-id.index (error files empty?)
- cat filename.sh.ojob-id.index (standard output files as expected?)
- tar czvf filename.tgz *
- mv filename.tgz \$HOME

Step 6 continued: sftp to get results

- exit
- Use sftp: from a new directory on your own computer of where you want the results to be. Open a shell and type ...
 - sftp username@login.cx1.hpc.ic.ac.uk
 - You will be asked for your standard cluster password
 - get filename.tgz
 - exit
- Your results are now all on your own computer
 - tar xzvf filename.tgz
- Your results are now complete uncompressed and ready for use. Now you need to write some R code to read in and analyze all those file.


For your exercises

- You'll be asked to adapt your code from yesterday to run on the cluster for a much bigger ecological community size
- You'll need to collect species abundance data as before and average over a large number of parallel simulations.



For your exercises

- You'll be asked to adapt your code from yesterday to run on the cluster for a much bigger ecological community size
- You'll need to collect species abundance data as before and average over a large number of parallel simulations.
- Use a "burn in" period and check the species abundance distribution periodically. You should plot species richness against time and make a conservative judgment, but for neutral theory $8 * \text{metacommunity size}$ complete turnovers of the community is a good rule of thumb.



DO NOT ...

- Use the cluster without knowing memory and time requirements
- Run jobs on the login node
- Try to use cx2 or ax4 parts of the cluster
- Output data to the hard disk regularly
- Use the same random seed for your simulations
- Copy and paste your shell script
- Leave your results in \$TMPDIR
- Waste too much of your own time optimizing your code
- Run code on the cluster that hasn't been tested locally first

DO ...

- Use the cluster for jobs that take a long time locally.
 - Optimize your code if there's going to be a huge benefit
 - Run repeat readings and different parameters as separate jobs.
 - Run jobs that take between 30 mins and 3 days to execute.
 - Write your shell script on the cluster itself.
 - Make your code output each result in a differently named file.
 - Check periodically that all is well on the cluster
 - Be ambitious – you can do loads of great stuff with a cluster.
- imperial.ac.uk/admin-services/ict/self-service/research-support/rcs
wiki.imperial.ac.uk/display/HPC/High+Performance+Computing